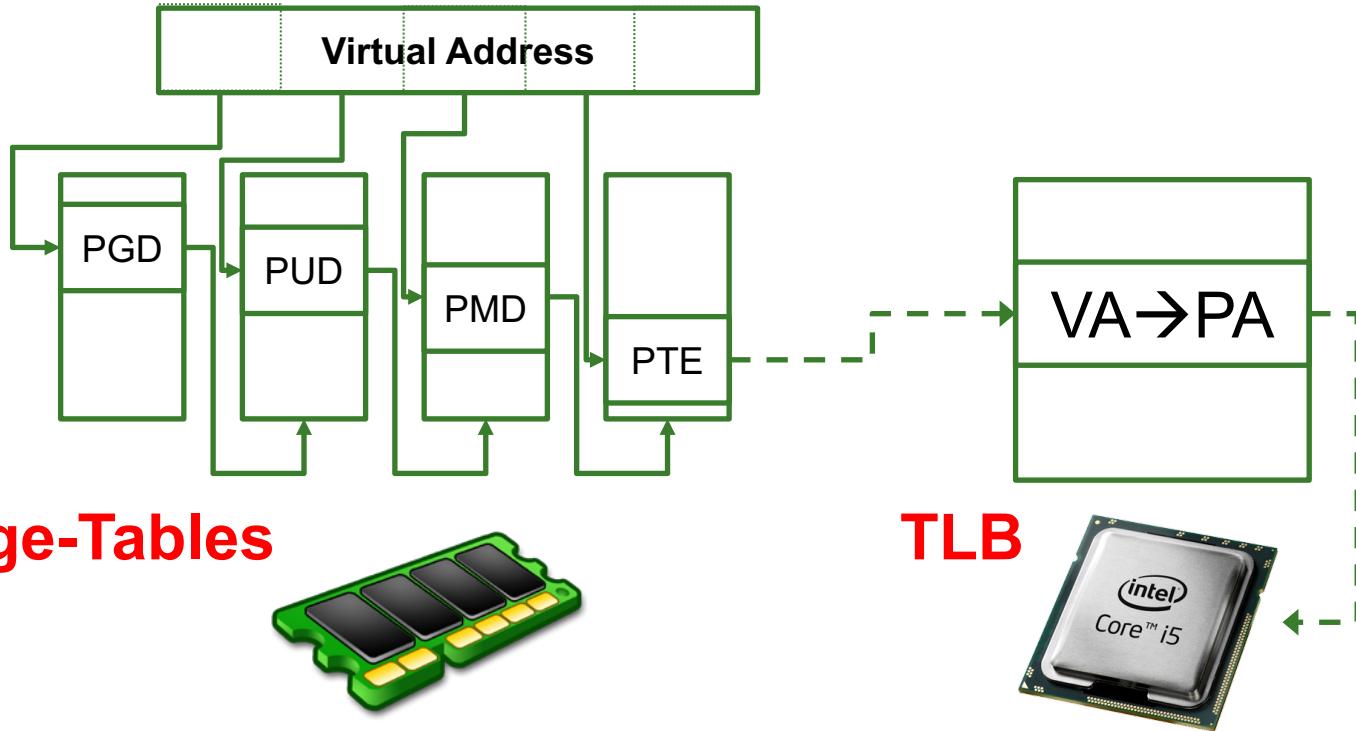


Optimizing the TLB Shootdown Algorithm with Page Access Tracking

Nadav Amit – VMware Research Group

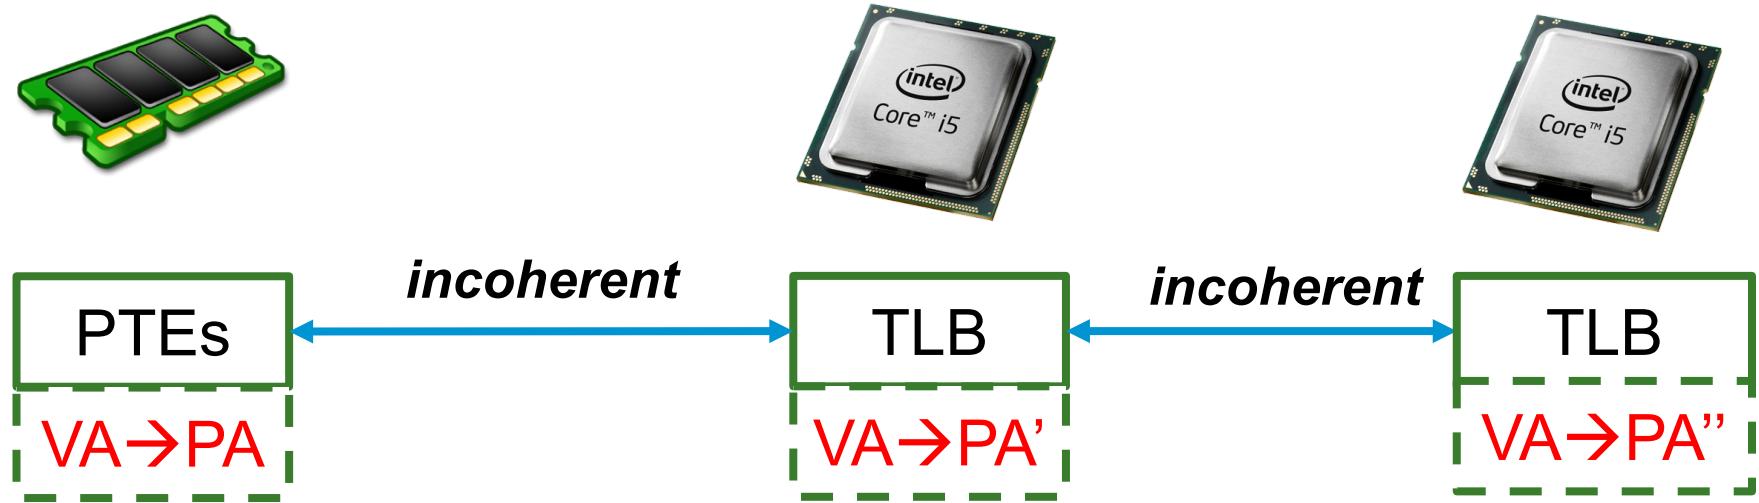


Translation Lookaside Buffer (TLB)



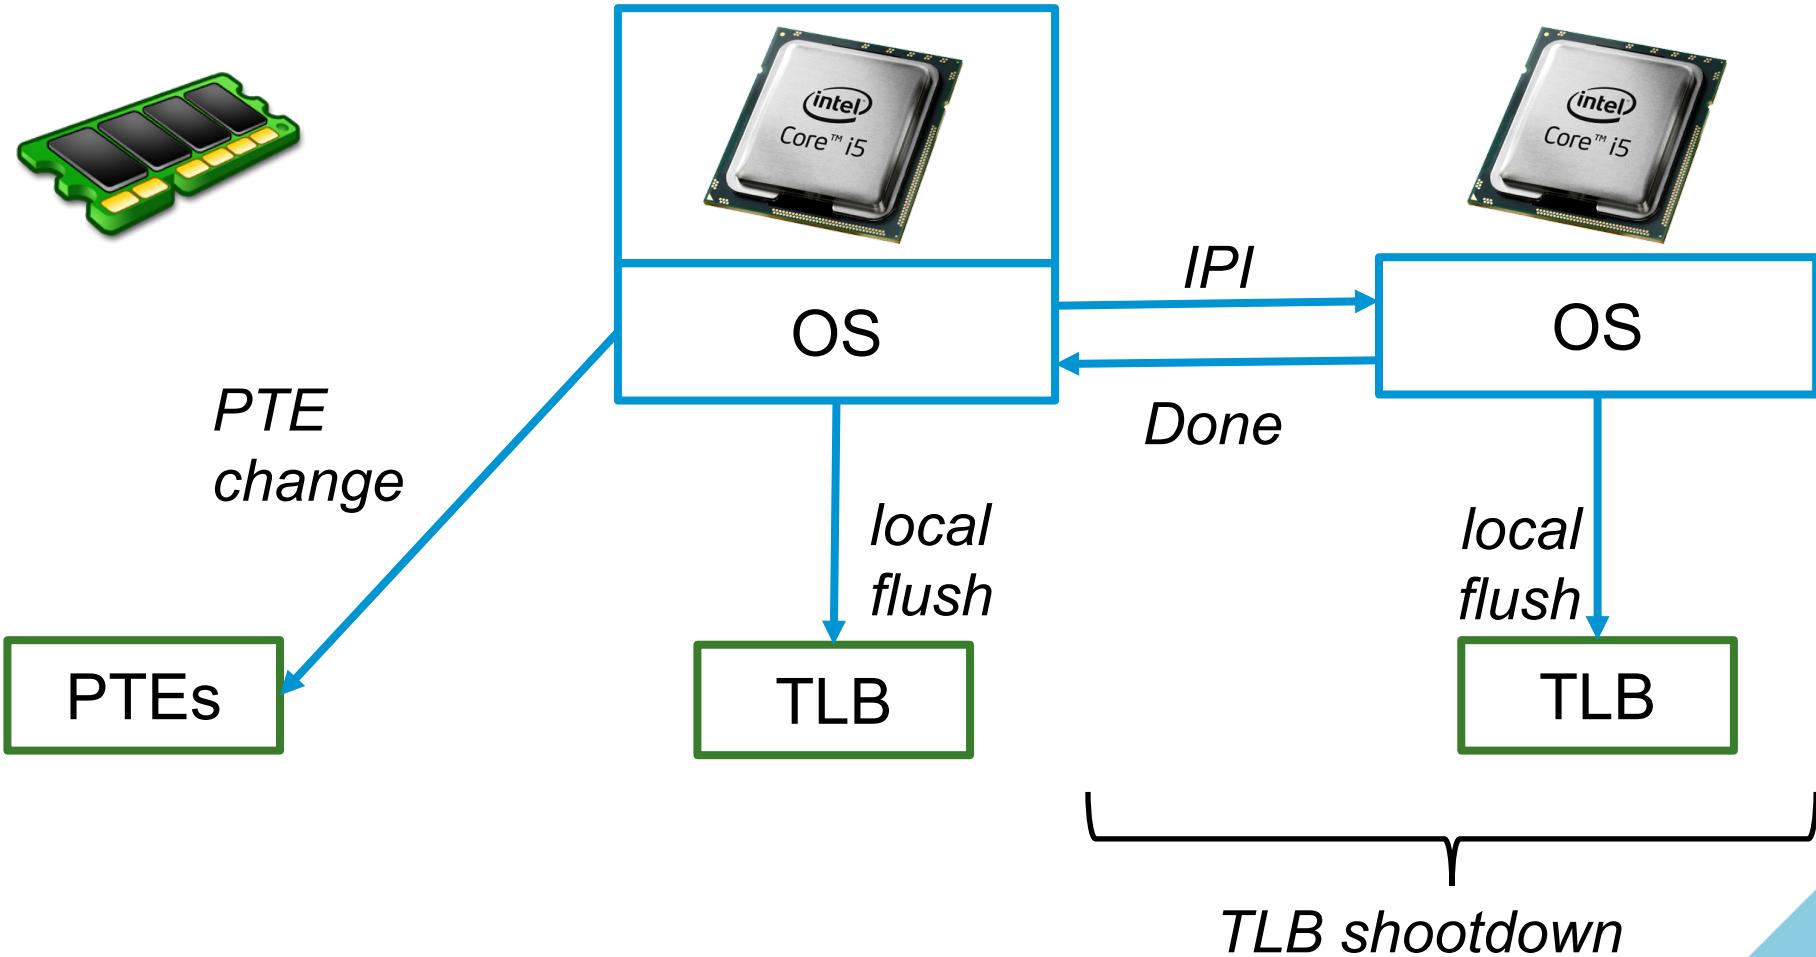
- TLB = cache for virtual to physical address translations

TLB Coherency



- Hardware does not maintain TLBs coherent
- The problem is left for software (OS)

Local TLB Flushes and Remote TLB Shootdowns



When do TLB Flushes Occur?

- Application initiated
 - munmap()
 - Copy-on-write
 - msync()
 - mprotect()
 - madvise()
 - migrate_pages()
- OS initiated
 - NUMA migrations
 - Memory compaction
 - Memory deduplication
 - Memory reclamation
 - Memory balloon
 - Background dirty-pages flush
- Faster storage → TLB overheads are more apparent

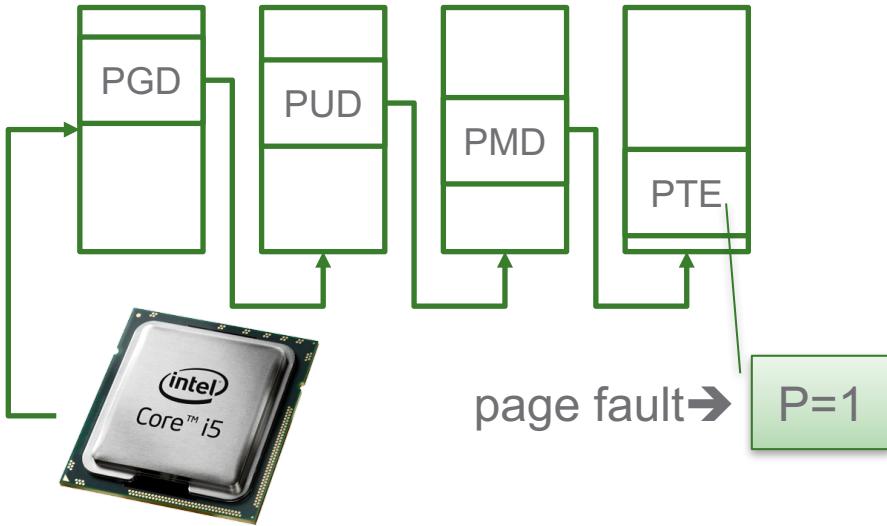
Existing Solutions

- Hardware [Teller'90, Villavieja'11, Li'13]
- Software (commodity OSes)
 - Batching [Uhlig'05]
 - Limit flushes to cores that use the address-space
 - Trade-off between full and individual PTE flushes
- Software (academic)
 - Explicit software control [Boyd-Wickizer'10, Tene'11]
 - Replicated paging hierarchy [Clements'13, Gerofi'13]

Replicated Paging Hierarchy

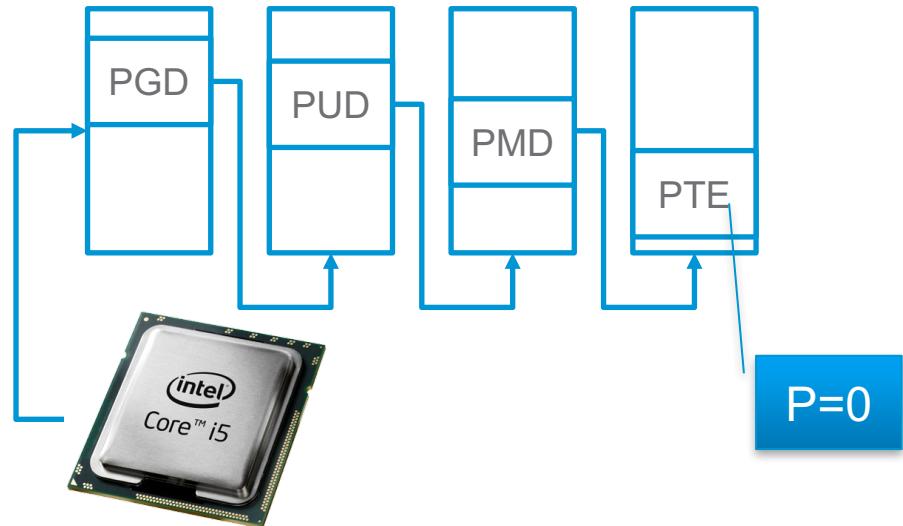
[Clements'13, Gerofi'13]

Page-Tables



page fault → P=1

Page-Tables



P=0

- Page-fault on each CPU that accesses a PTE
- Memory overheads
- Runtime overheads: managing multiple tables

Insight: Use PTE Access-Bit

PTE:

page frame number	A	permissions
-------------------	---	-------------

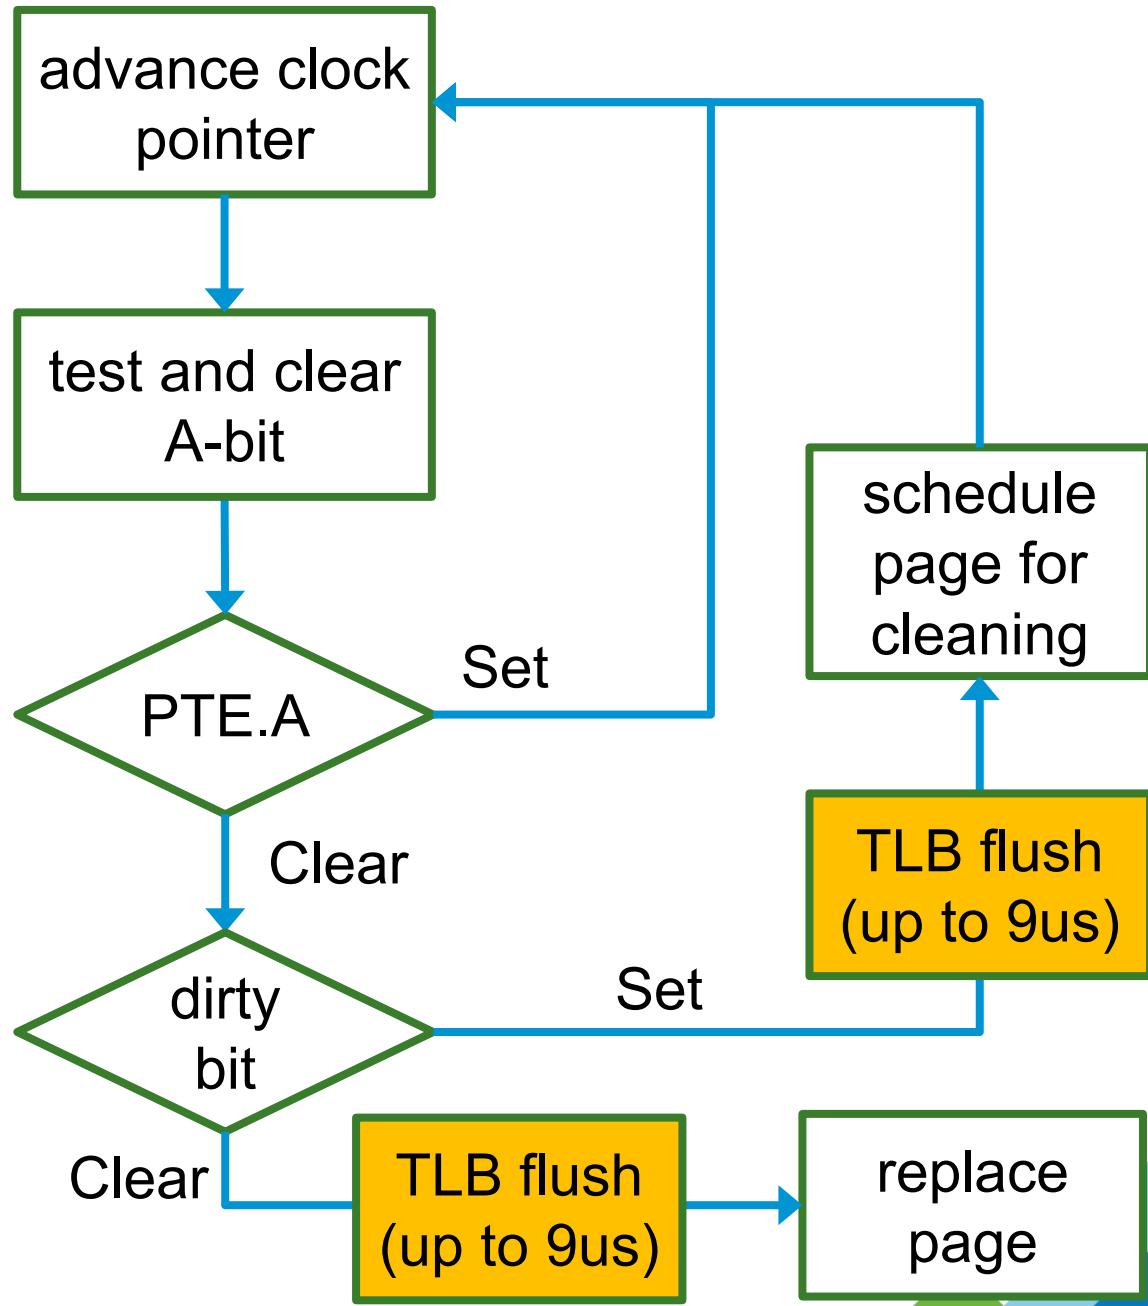
- Set by hardware, cleared by software
- Used for OS **memory reclamation decisions**
 - Set when a page is **accessed**
 - “These flags are provided ... to manage the transfer of pages ... into and out of physical memory.” (Intel SDM)
- Insight: **can be used for TLB invalidation decisions**
 - Set when a PTE is **cached**
 - “Whenever the processor uses a PTE as part of address translation, it sets the accessed flag...” (Intel SDM)

Our System

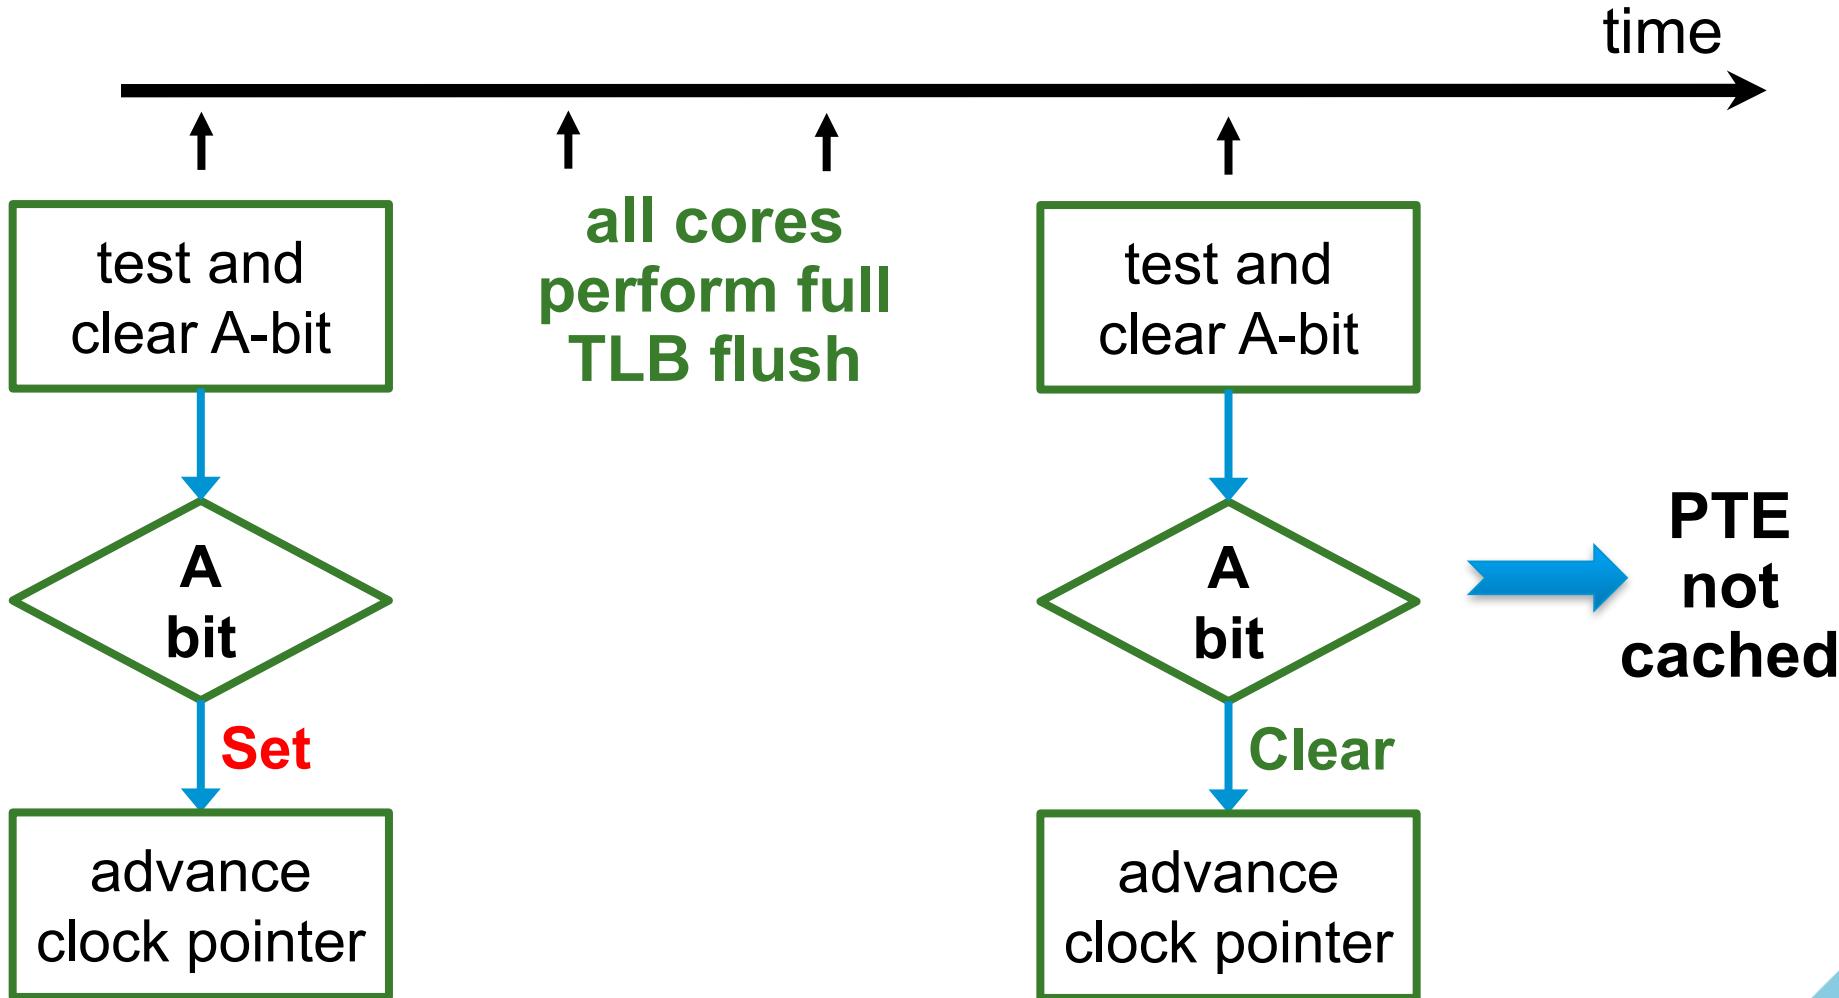
- Flush decisions based on PTE access-bit
- Software solution (x86)
 - Exploiting the full potential requires simple hardware changes
- Prevent common unnecessary TLB shootdowns
 - Long-lived idle mappings
 - Short-lived private mappings
- Some false positives
 - Unnecessary flushes

Long-Lived Idle Mappings

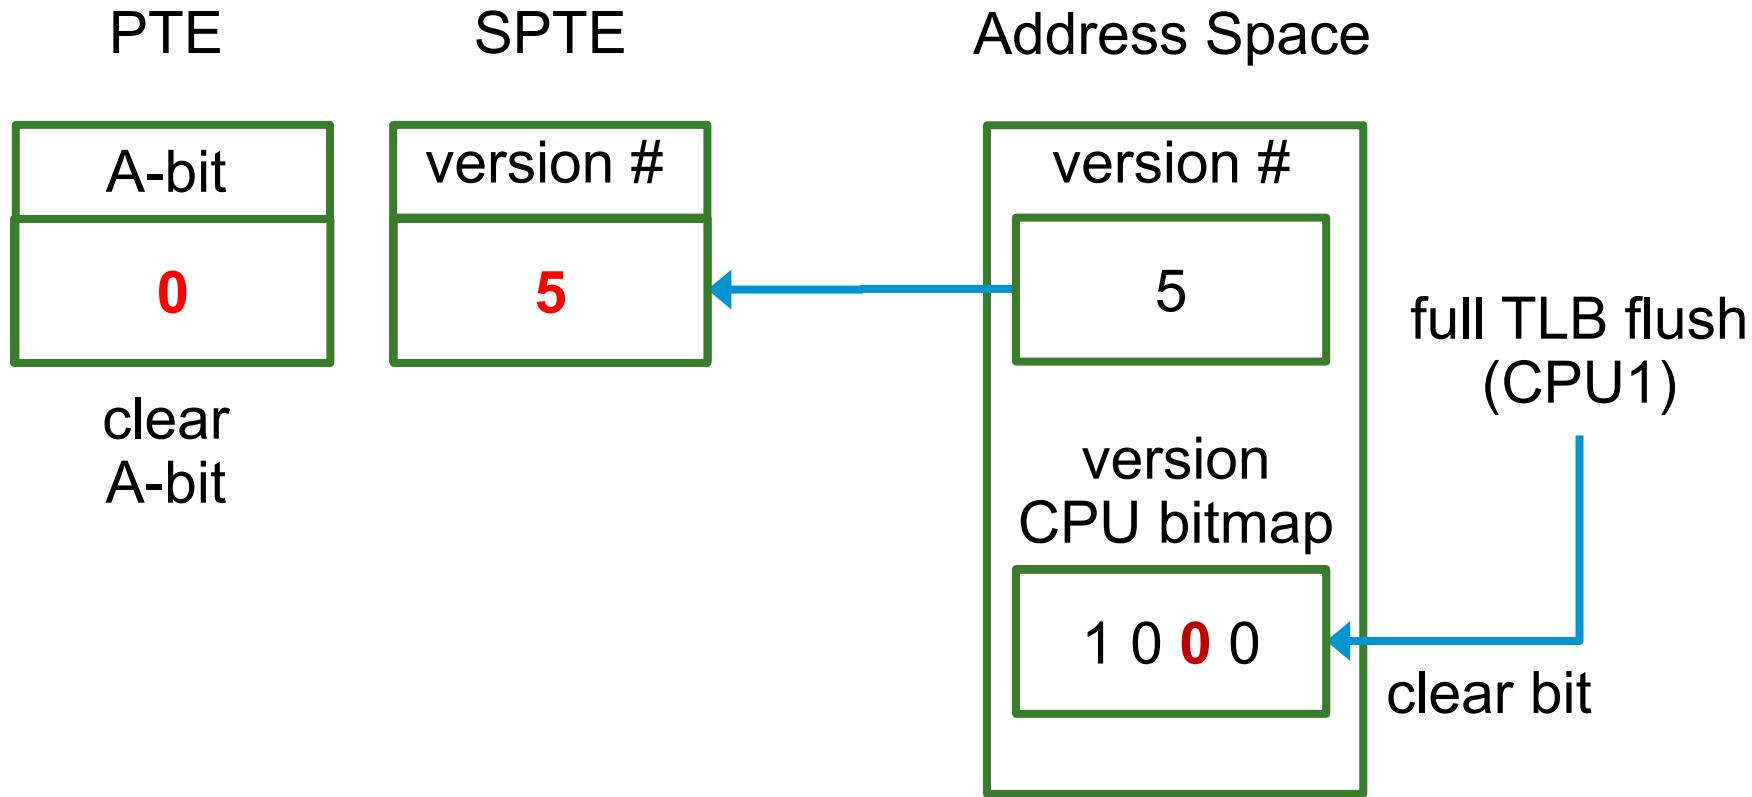
CLOCK Algorithm
[Carr and Hennessy '81]



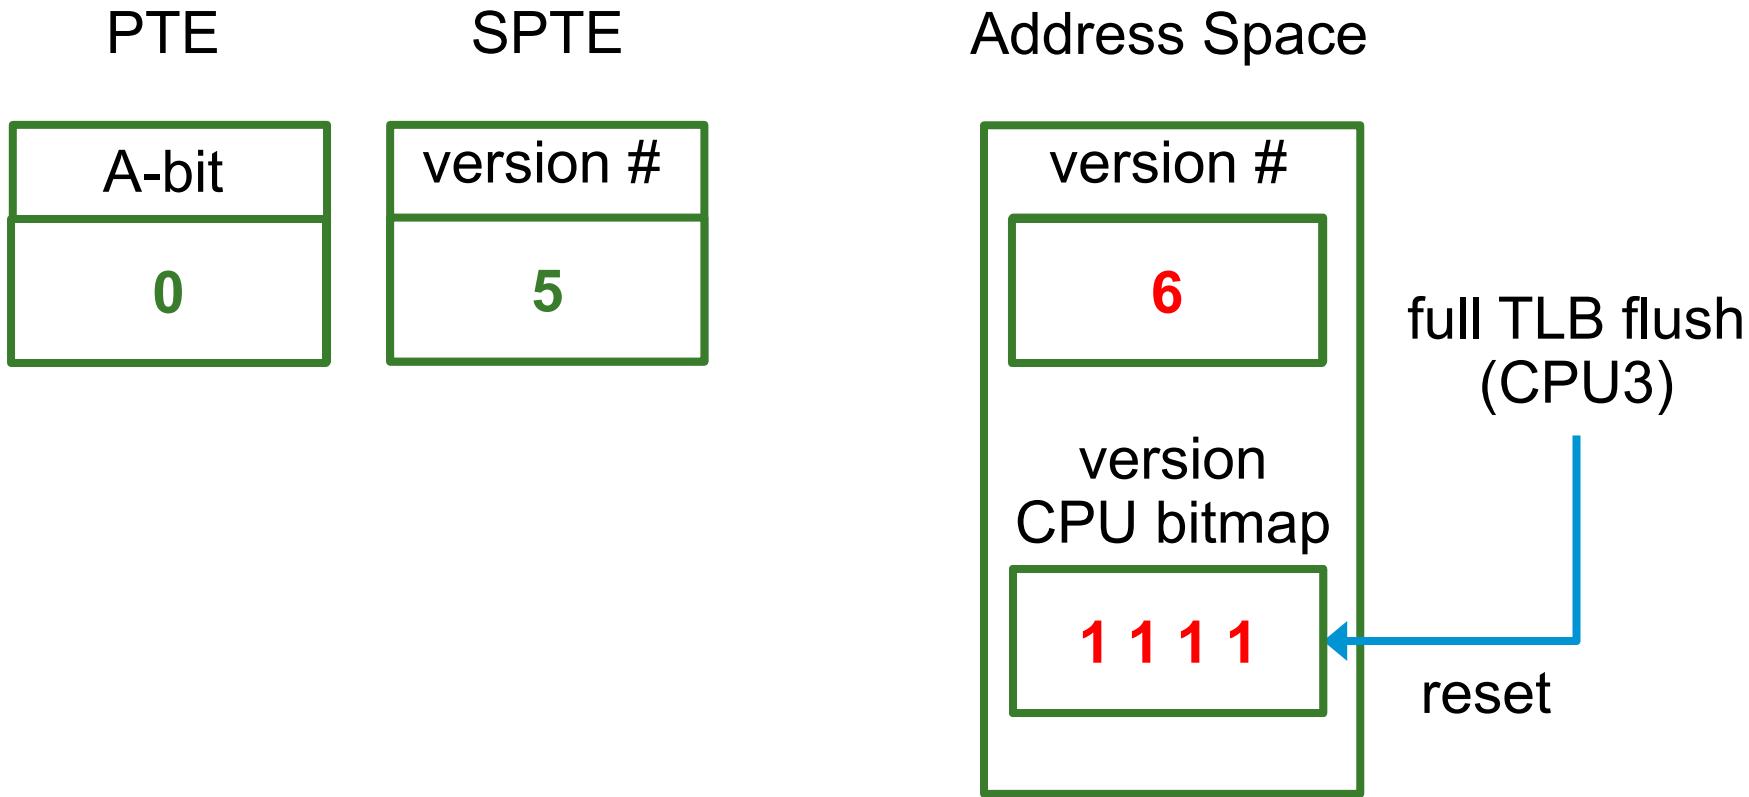
Avoiding Flush of Long Lived PTEs



TLB Version Tracking (1)

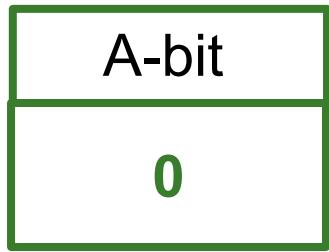


TLB Version Tracking (2)



TLB Version Tracking (3)

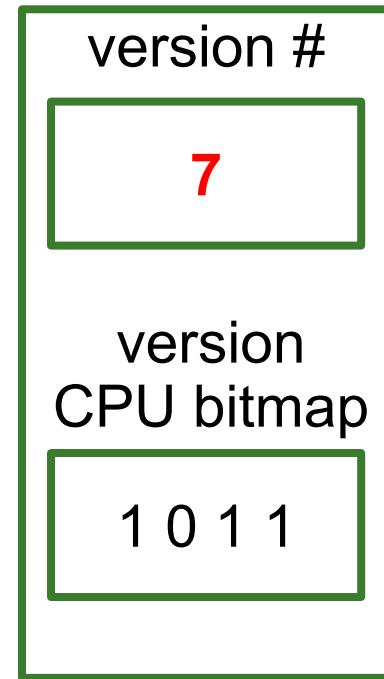
PTE



SPTE



Address Space

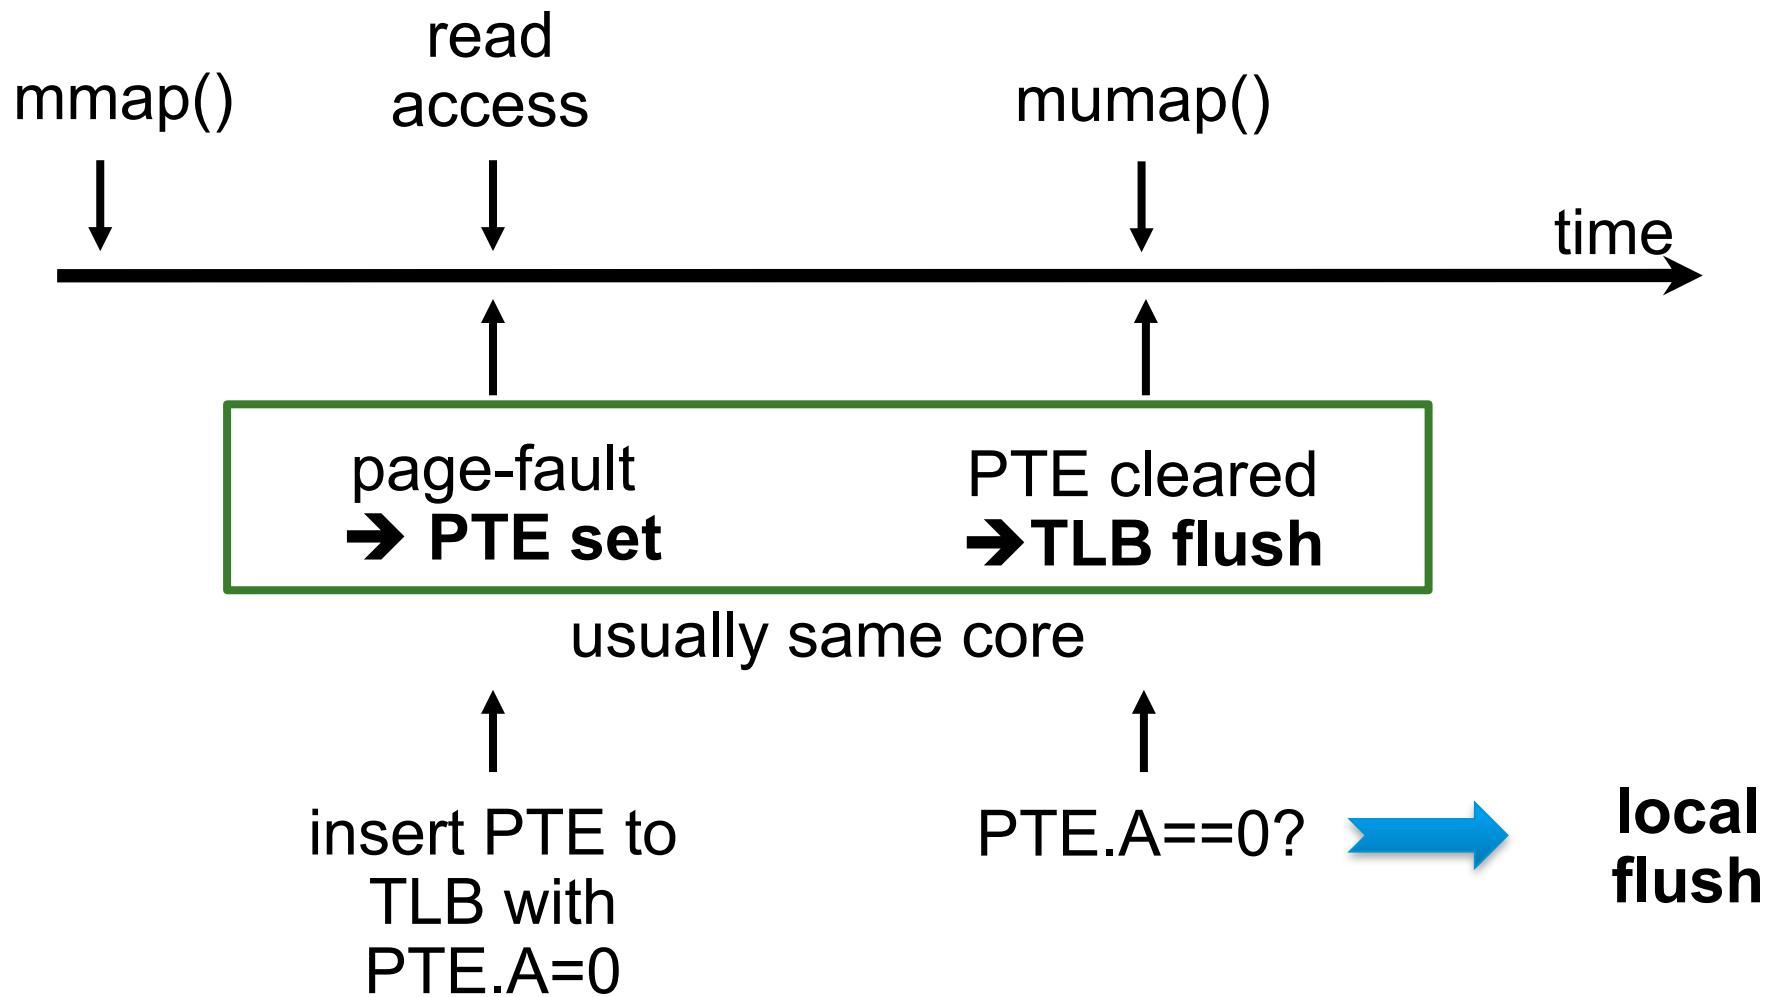


PTE unmap and flush:

If $PTE.A == 0$ and
 $SPTE.ver + 1 < AS.ver$

Then avoid TLB flush

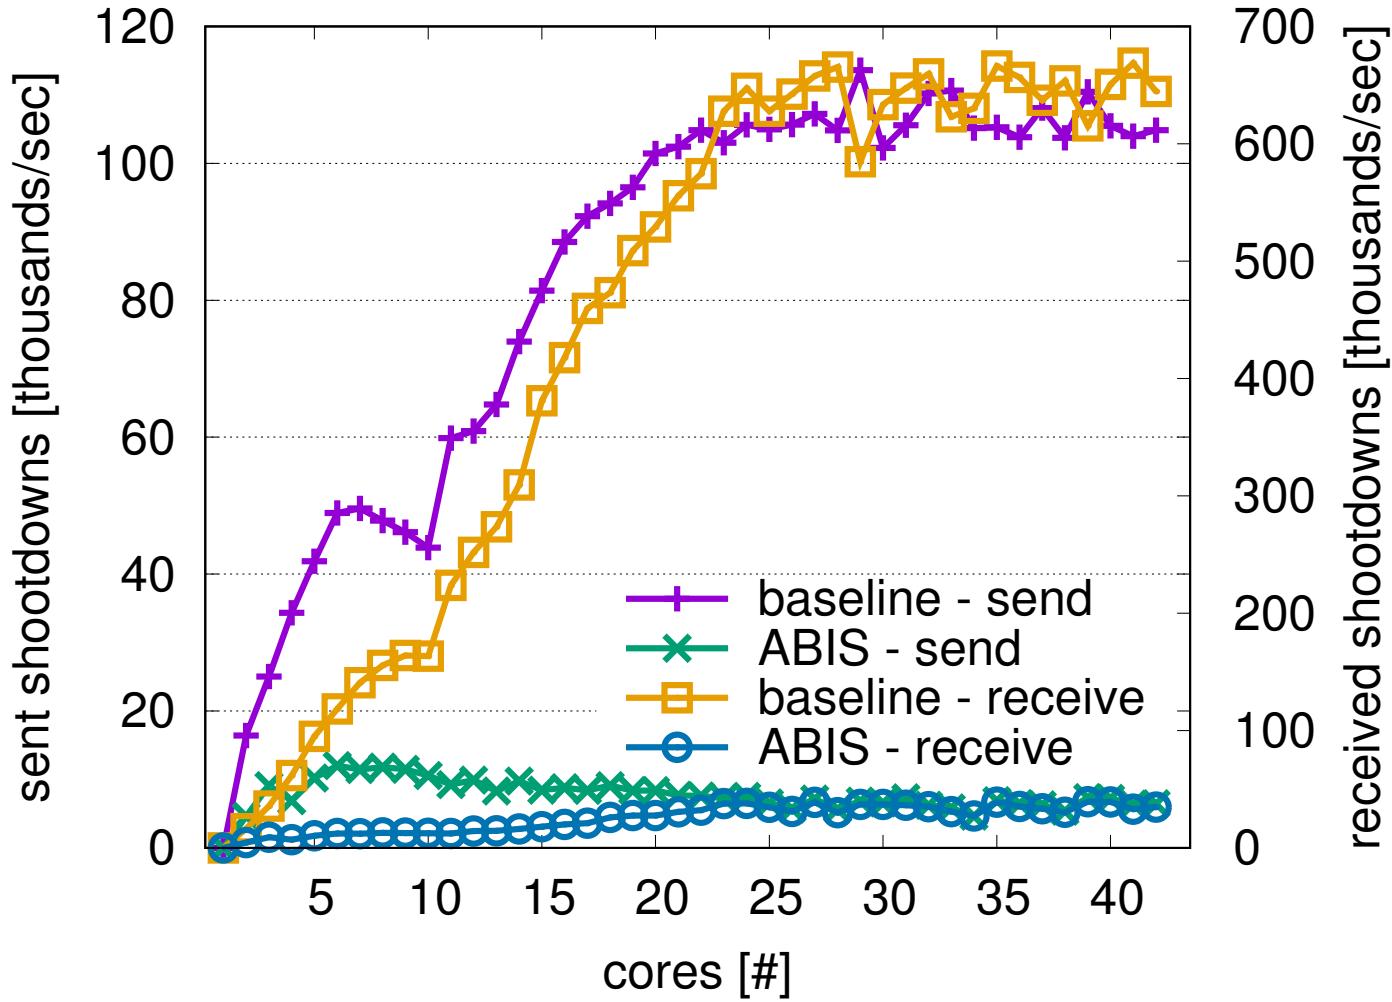
Short Lived Private Mappings



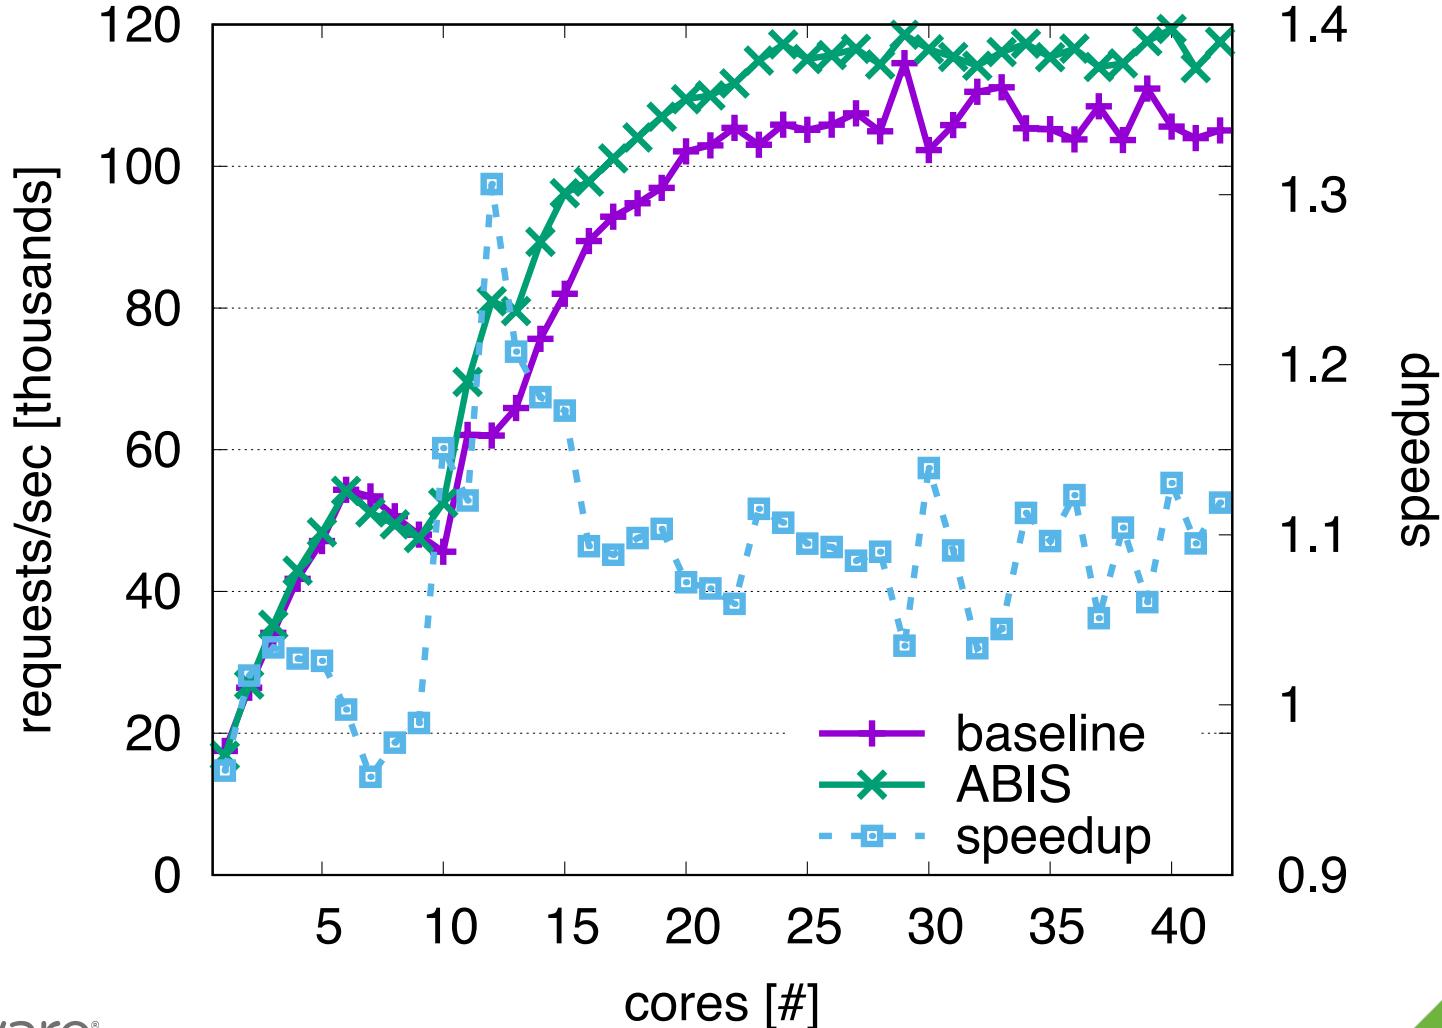
Evaluations

- Prototype based on Linux 4.5
- Baseline configured to avoid shootdown cost
 - Linux version that uses TLB flushes batching
 - Using efficient multicast IPI delivery
- 48-cores, 2-socket server
- Our system denoted as ABIS:
Access-Based Invalidation System

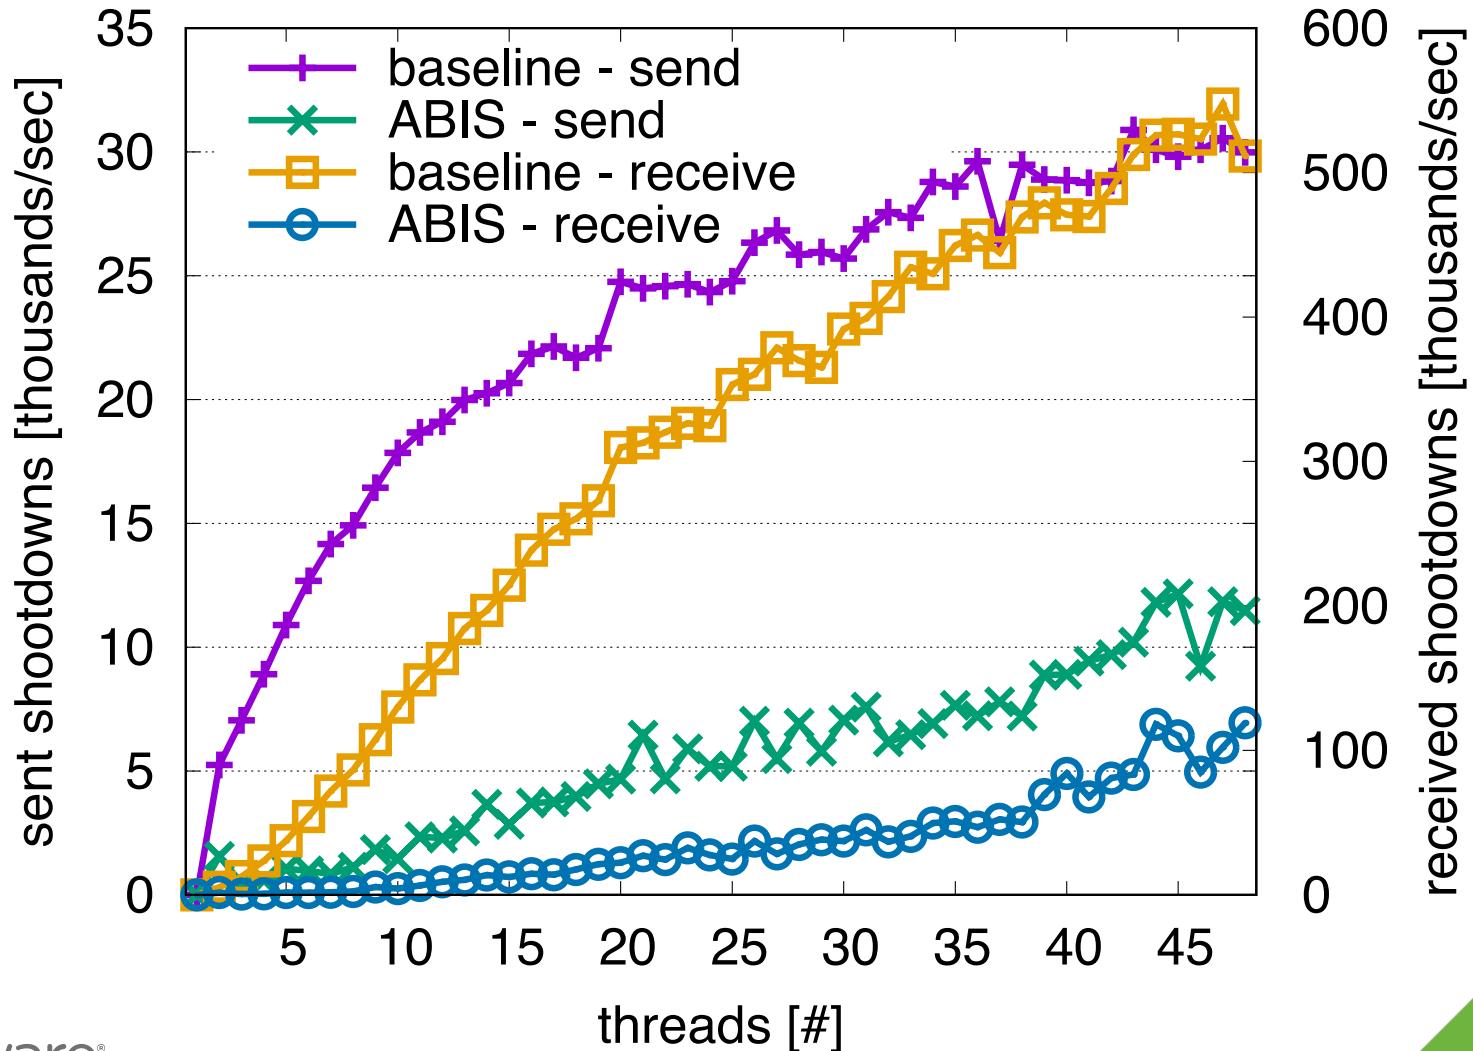
Apache TLB Shootdowns (Short-Lived Private Mappings)



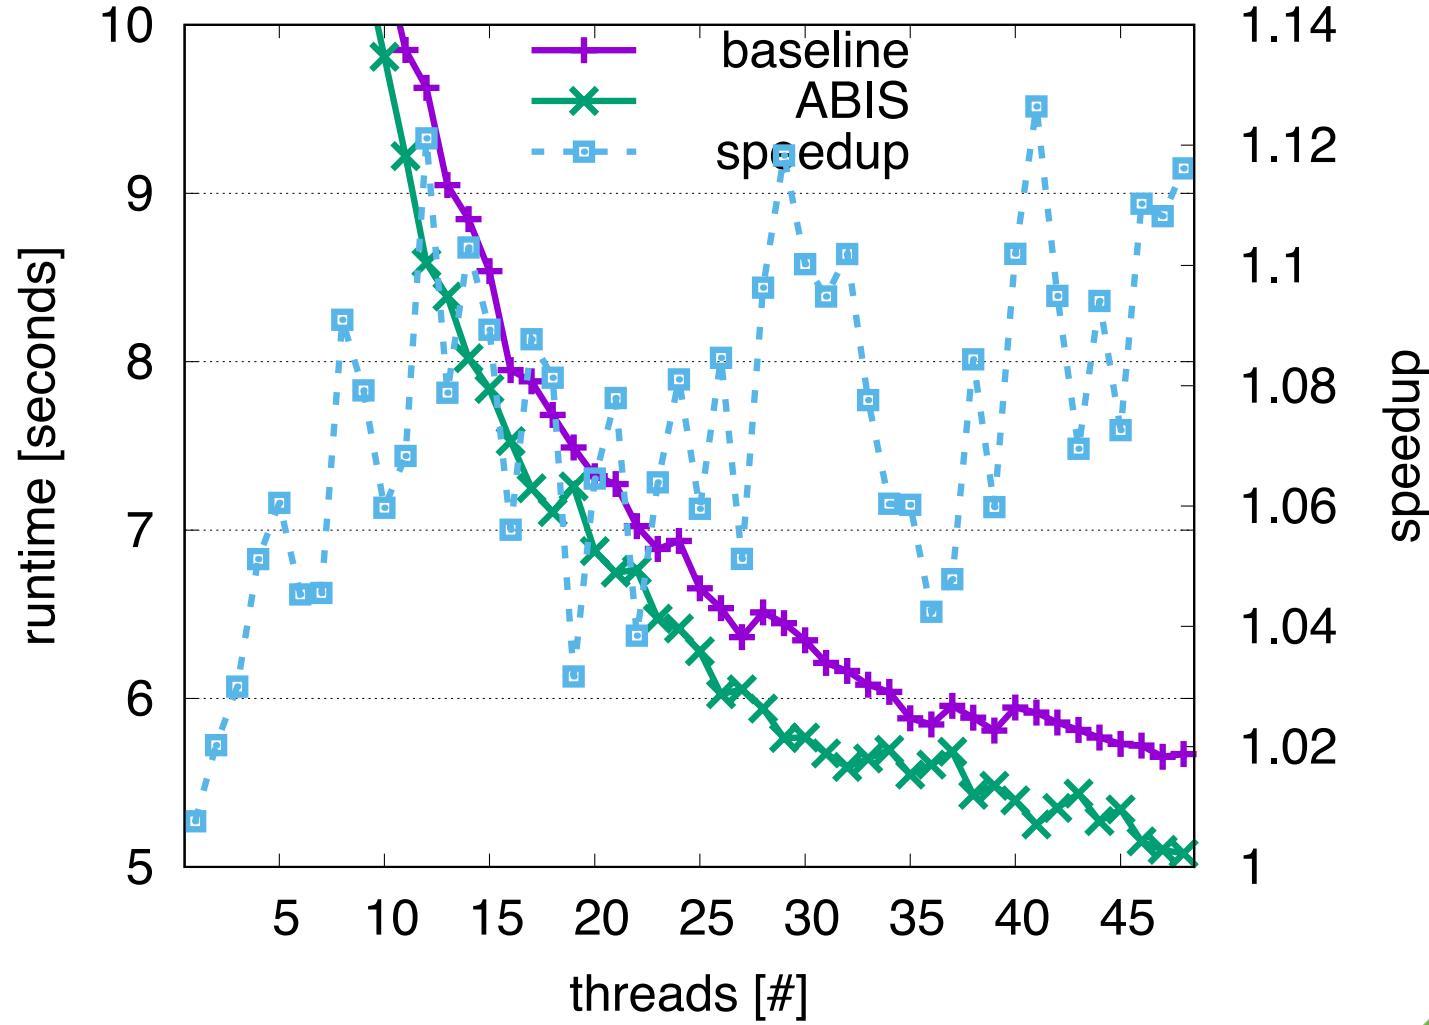
Apache Performance



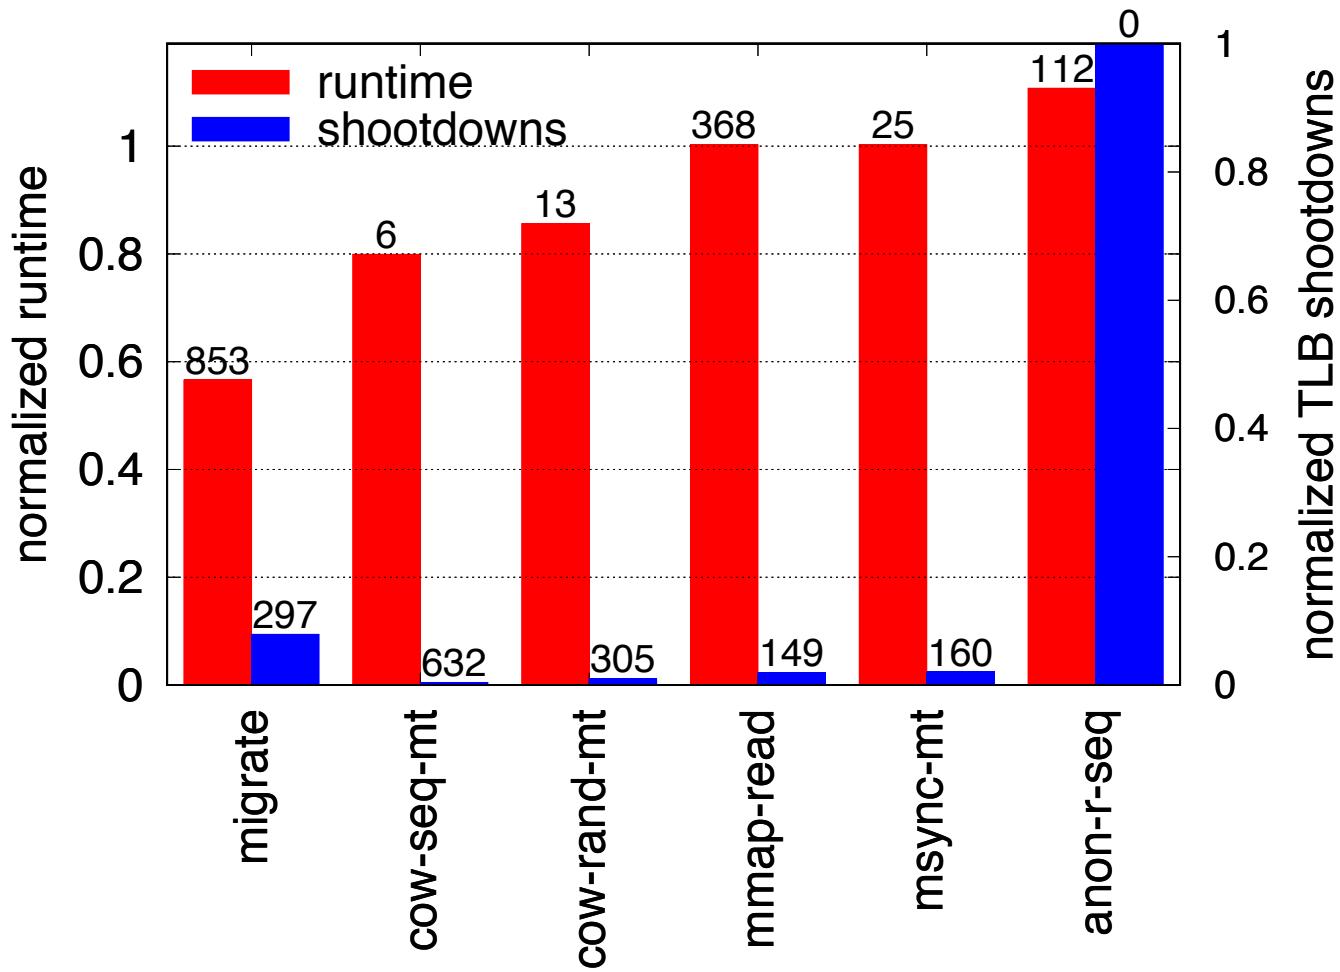
PBZIP2 – TLB shootdowns (Long-Lived Idle Mappings)



PBZIP2 Performance



Microbenchmarks: VMscale



Conclusions

- Access-bit tracking can often prevent most TLB shootdowns:
 - Long-lived idle PTEs
 - Short-lived private PTEs
- Exploit memory coherency to check if TLB is cached
- CPUs should allow more control over the TLB
 - Insertion of PTEs directly to the TLB