

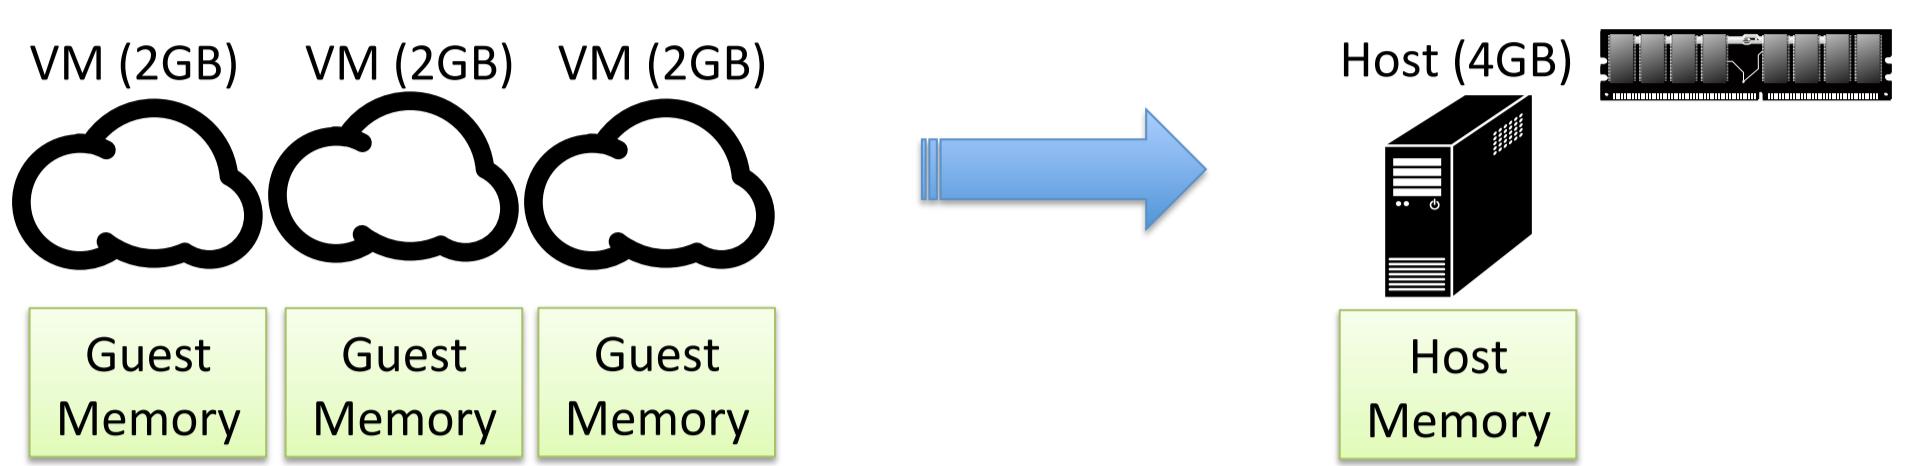
# A Memory Swapper for Virtualized Environments

Nadav Amit, Dan Tsafrir, Assaf Schuster



## Virtual Machines Memory Overcommitment

Memory overcommitment allows the hypervisor to oversubscribe memory to virtual machines.



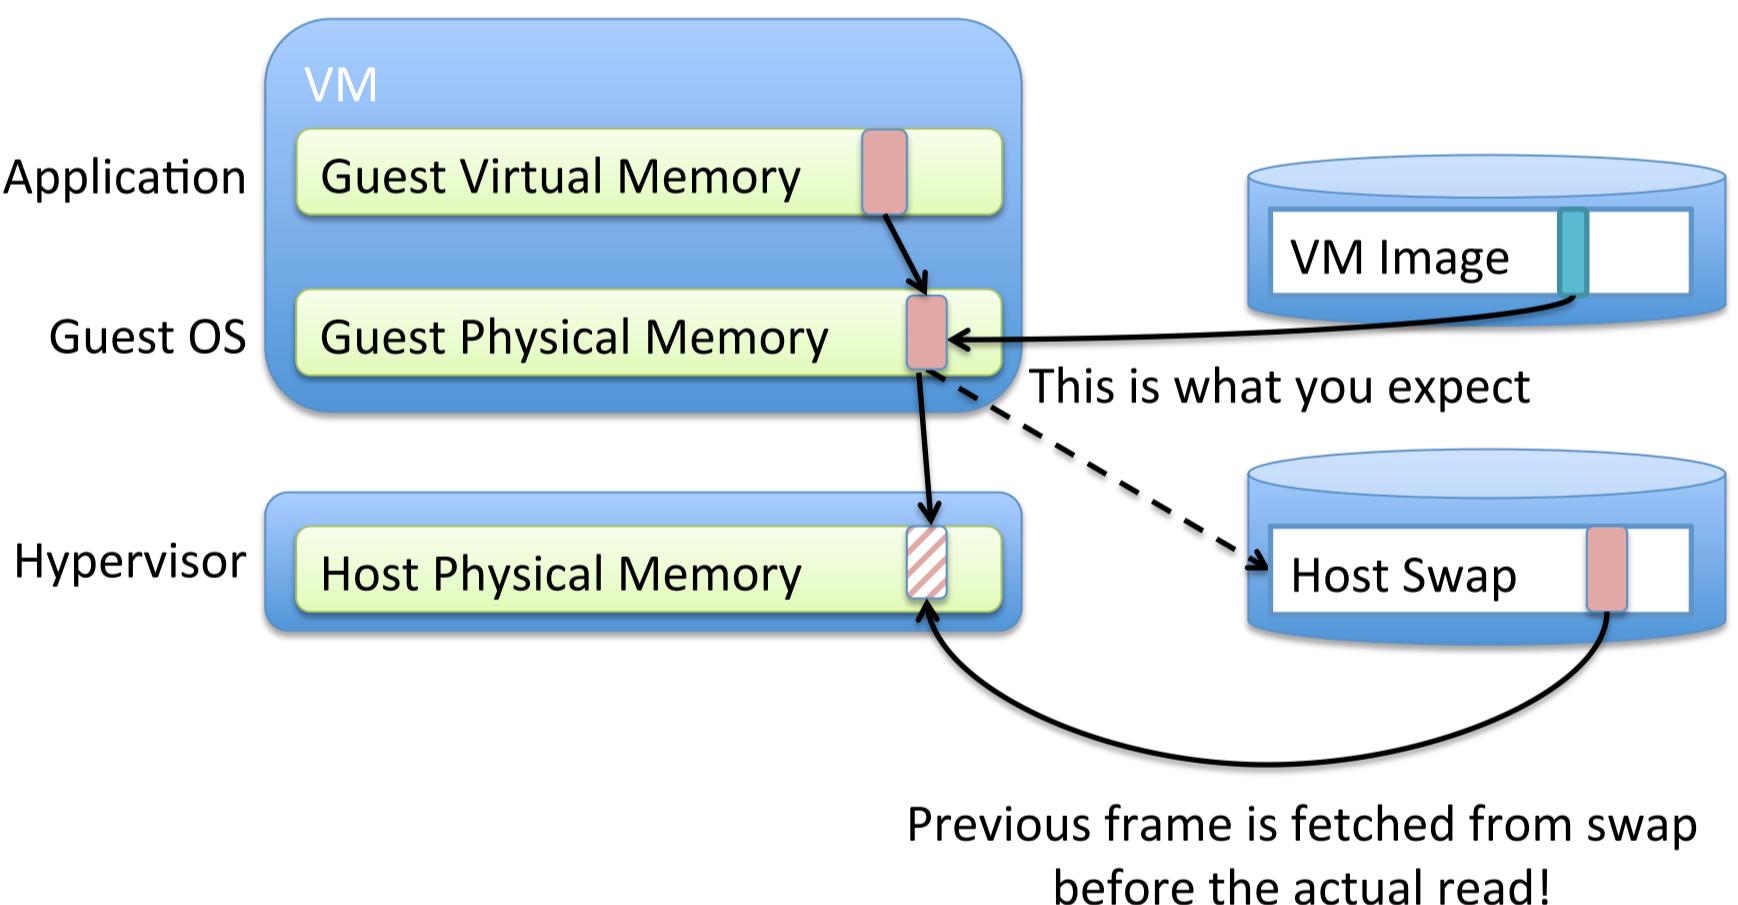
Two prevailing solutions:

Balloon – paravirtual, limited use, slow to respond

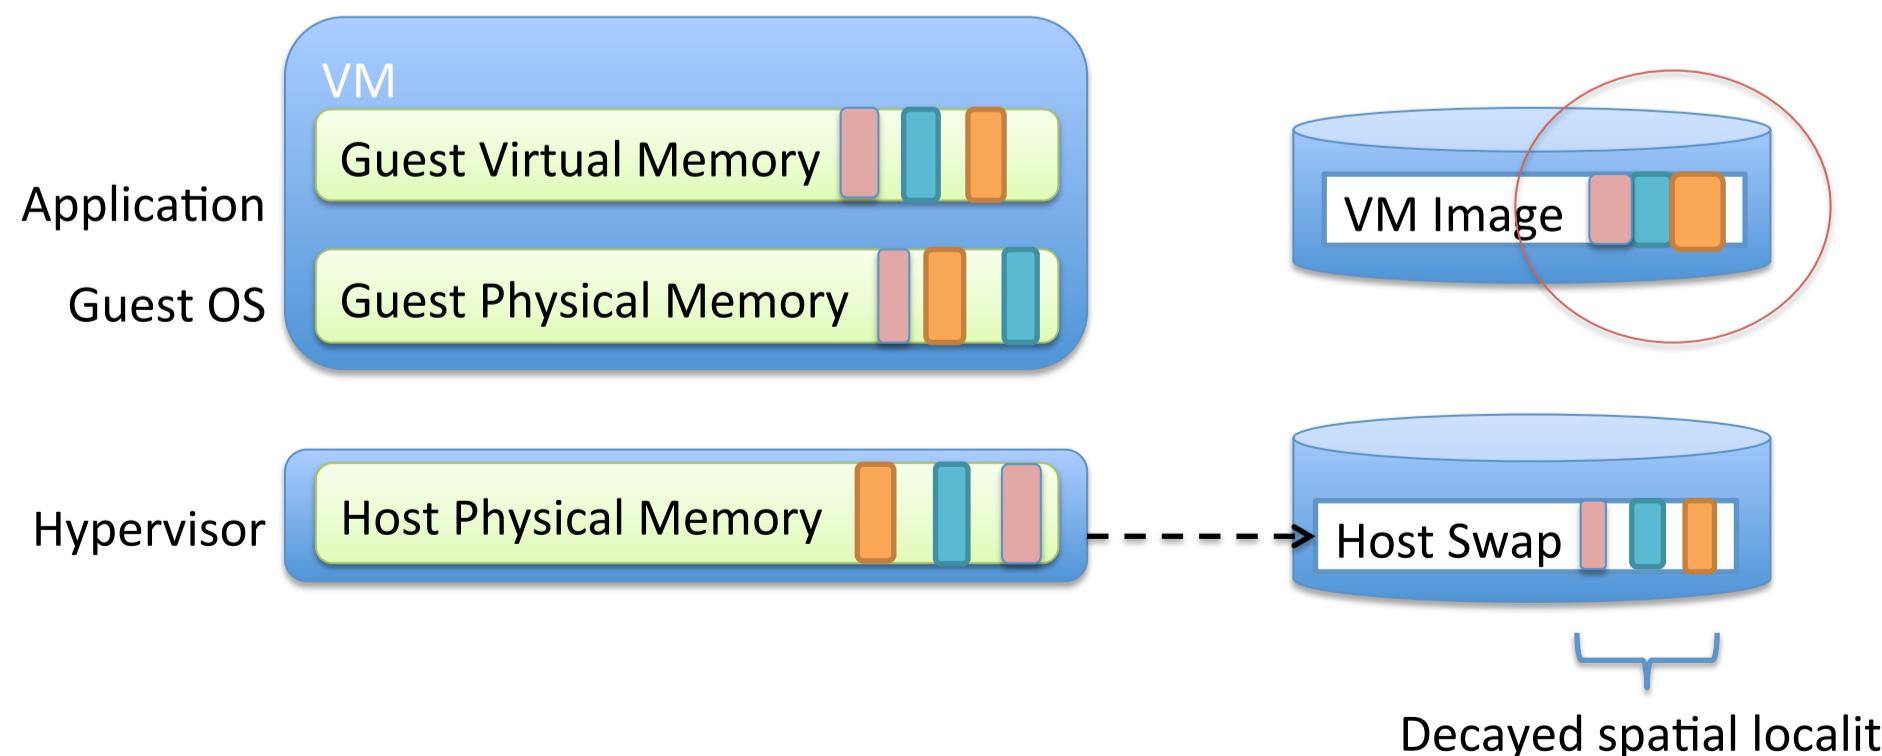
Host swapping – fully virtual, poor performance

## Analysis: Uncooperative Swap Overheads

1. Reads of stale memory which is being overwritten. The case of a VM disk read is illustrated.



2. Memory backed by VM disk is swapped out, marked anonymous and loses spatial locality.



3. False anonymity of guest memory. OS PFRA prioritize page-cache evictions.

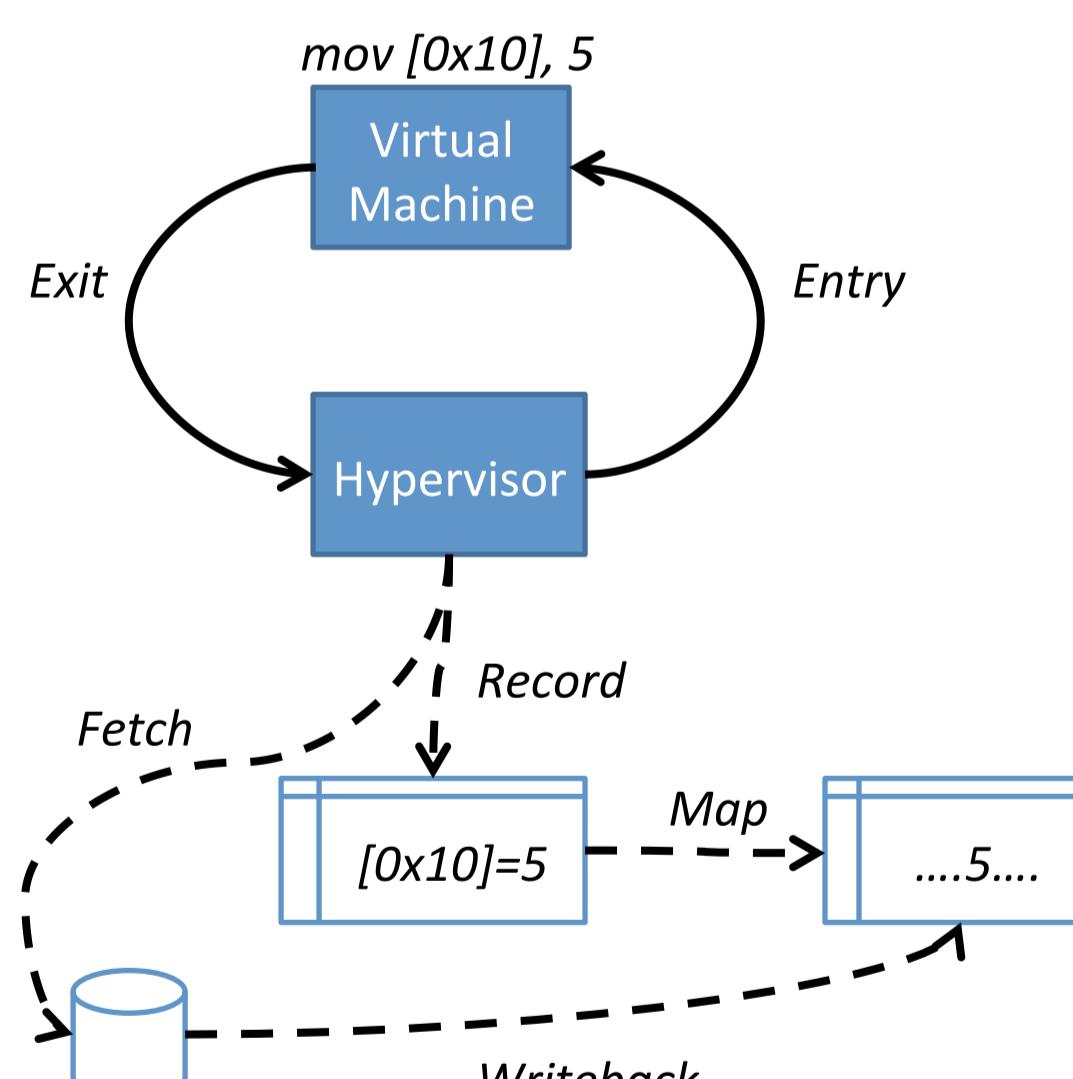
## Efficient Host Swapping

### Swap Mapper

- Clean pages which can be backed by the disk are mapped to disk blocks (new *mmap* variant).
- Mapped pages are discarded on eviction.
- Consistency is maintained on writes to blocks that back memory (unlike *mmap*).
- Prefetch only mapped memory.
- Use 4KB disk blocks.

### False Read Preventer

- Hypervisors cause false swap reads when the guest OS reallocates page frames (e.g., zeroing pages, COW).
- VSWAPPER eliminates these reads by emulating write instructions to non present memory and writing the result into write buffers.
- Full write buffer pages can be mapped to VM.
- Disk reads can be saved.



## Related Work

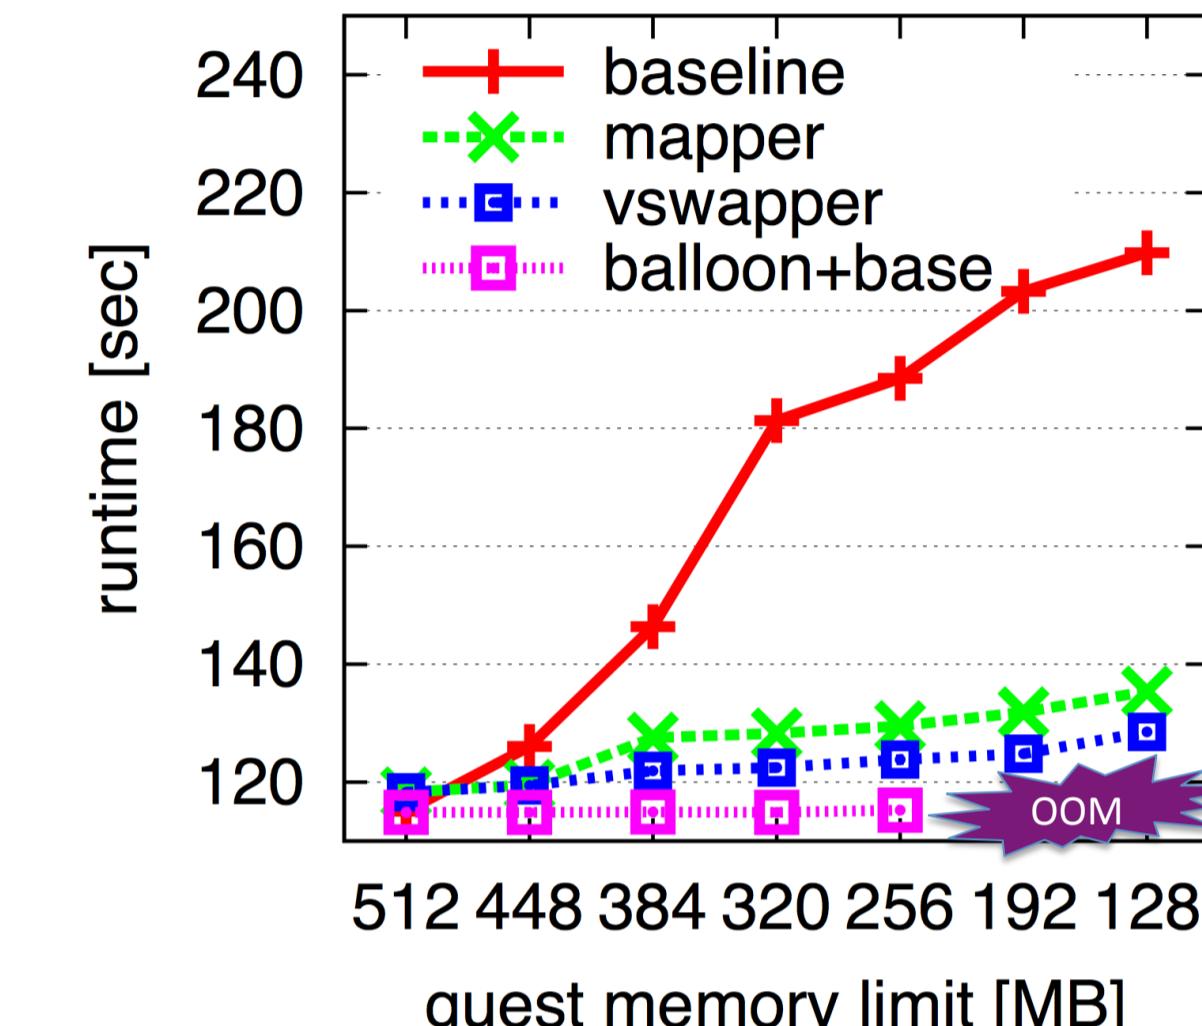
- Monitoring the buffer cache [Jones06]
- Copy-on-write disks [Bugnion97]
- VM memory access tracing [Lu07]
- Asynchronous page faults [Useche12]

## Contribution

1. Analysis of uncooperative swapping overheads
2. Swap Mapper to save memory-disk associations
3. Swap Preventer to avoid swap on reallocation
4. Significant improvement of common workloads

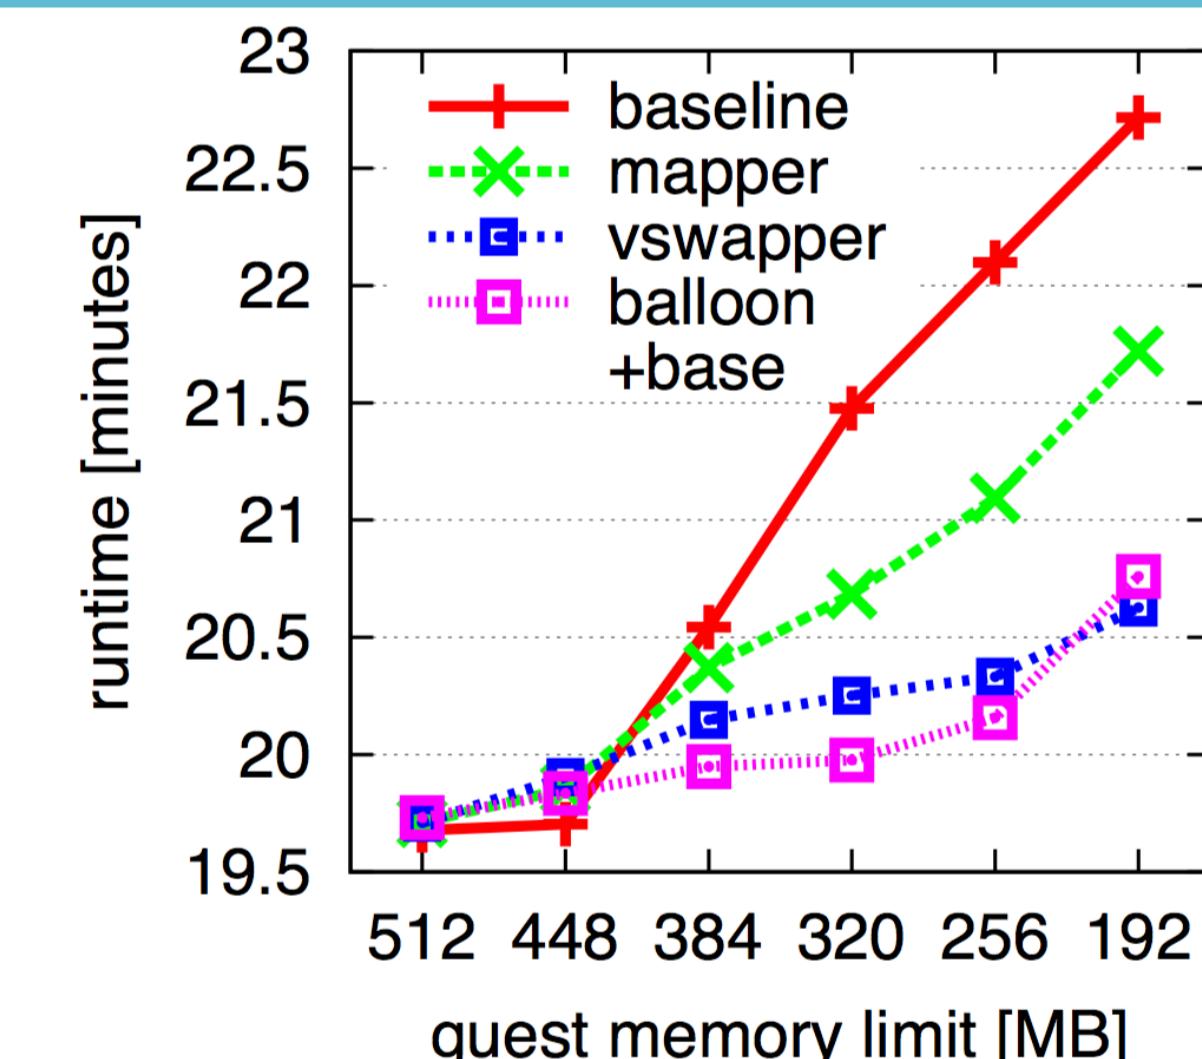
## Static Results: KVM, 512 MB guest, 1 VCPU

### PBZIP2 (Parallel BZIP2)



Avoiding stale reads and false anonymity improves performance by up to 1.63x.

### Kernbench (Linux Kernel Compilation)

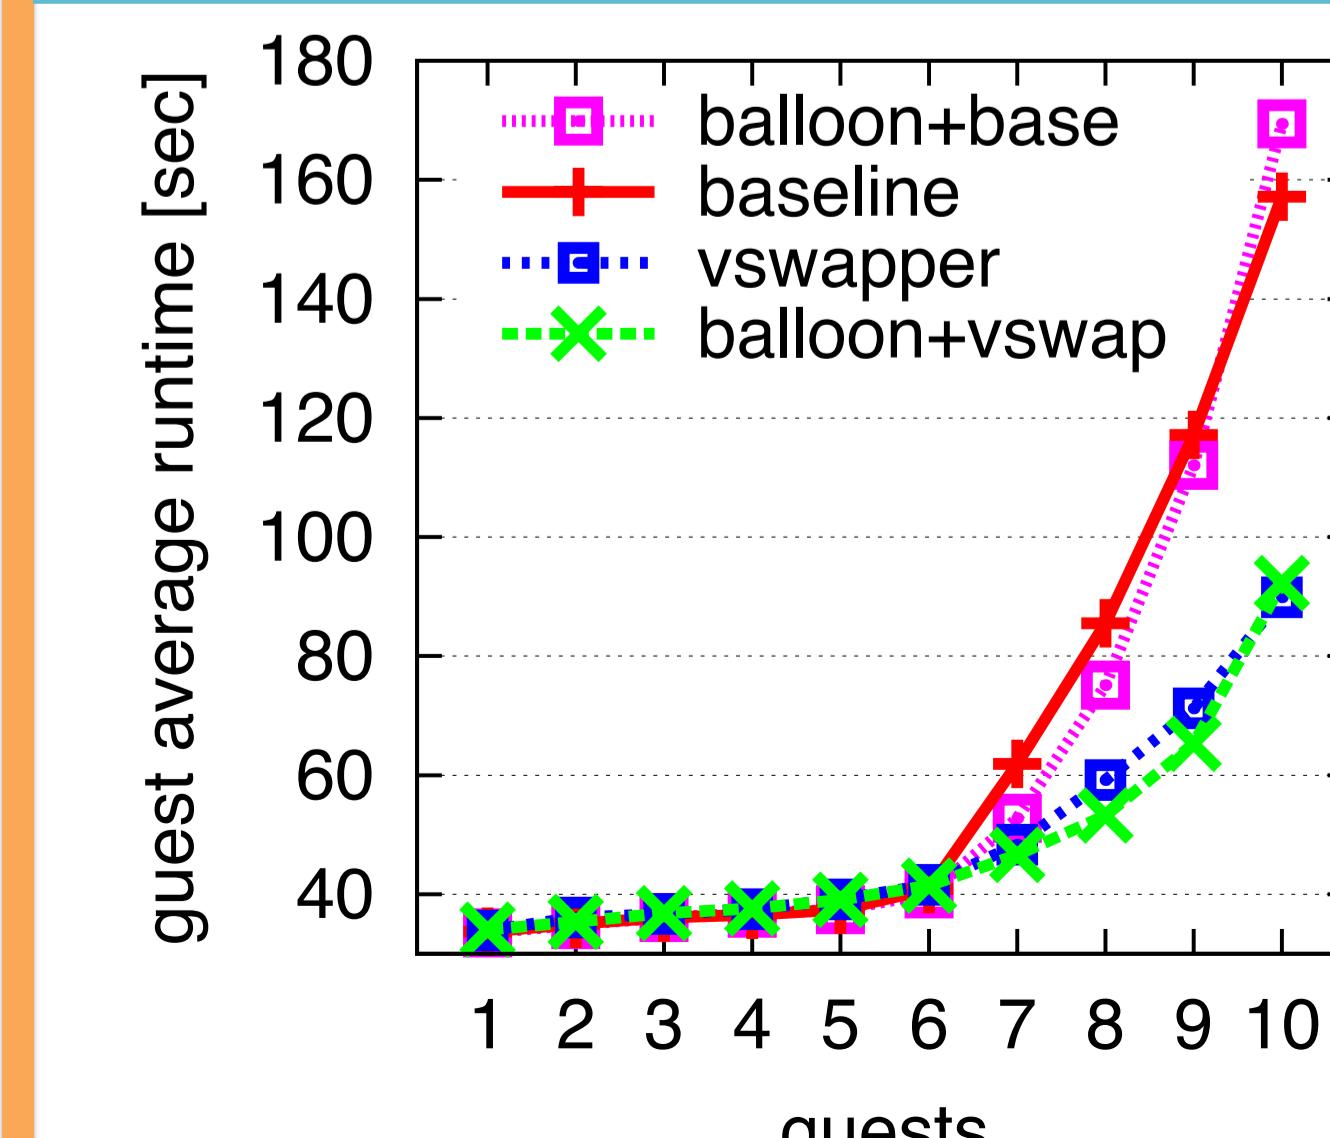


Up to 1.1x performance improvement.

Swap Preventer saves thousands of swap operations.

## Dynamic Results: 2GB guest – 8GB Host w/MoM

### MapReduce with multiple VMs



Balloon is less effective for dynamic workloads.

VSWAPPER delivers up to 2x improvement with ballooning